

COMP 448/548: Medical Image Analysis

Convolutional neural networks

Çiğdem Gündüz Demir

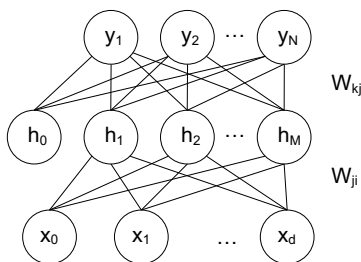
cgunduz@ku.edu.tr

1

Multilayer perceptrons

(revisited)

- Contain hidden layers



Hidden units h_j 's can be viewed as new "features" obtained by combining x_i 's

A deeper architecture with nonlinear activations is more expressive than a shallow one

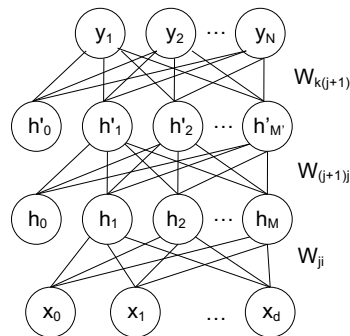
In this network

- Each hidden unit computes its net activation
$$net_j^t = \sum_i x_i^t W_{ji}$$
- Each hidden unit emits an output that is a nonlinear function (e.g., sigmoid, ReLU) of its activation
$$h_j^t = \text{non-linear-function}(net_j^t)$$
- Each output unit computes its net activation
$$net_k^t = \sum_j h_j^t W_{kj}$$
- Each output unit emits an output (using a linear, a sigmoid or a softmax function)
$$y_k^t = \text{output-function}(net_k^t)$$

2

More hidden layers

(revisited)



$$\frac{\partial \text{loss}^t(W)}{\partial W_{ji}} = \frac{\partial \text{loss}^t(W)}{\partial \text{net}_j^t} \cdot \frac{\partial \text{net}_j^t}{\partial W_{ji}}$$

$$\frac{\partial \text{loss}^t(W)}{\partial W_{ji}} = \delta_j^t \cdot x_i^t$$

$$\delta_j^t = \sum_{(j+1)} \frac{\partial \text{loss}^t(W)}{\partial \text{net}_{(j+1)}^t} \cdot \frac{\partial \text{net}_{(j+1)}^t}{\partial \text{net}_j^t}$$

$$\delta_j^t = \left[\sum_{(j+1)} \delta_{(j+1)}^t \cdot W_{(j+1)j} \right] \cdot \sigma'(\text{net}_j)$$

δ_j may vanish after repeated multiplication. This makes deep architectures hard to train (when initial weights are not “good” enough)

Approaches for alleviating underfitting and overfitting problems

- Better network designs: [Sparse connections](#), [weight sharing](#), [convolutional nets](#), long/short skip connections, [activation functions](#), ...
- Better network training: Regularization, loss function definitions, larger datasets, data augmentation, ...
- Previously, [layerwise pretraining](#) ([restricted Boltzmann machines](#), [autoencoders](#))

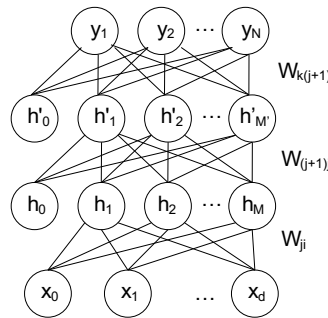
3

Layerwise pretraining

4

Vanishing gradient problem

- This makes hard to train deep architectures (with many hidden layers) by backpropagation
- When the initial weights are good enough, backpropagation works well
- Layerwise pretraining
 - Restricted Boltzmann machines
 - Autoencoders



$$\frac{\partial \text{loss}^t(W)}{\partial W_{ji}} = \frac{\partial \text{loss}^t(W)}{\partial \text{net}_j^t} \cdot \frac{\partial \text{net}_j^t}{\partial W_{ji}}$$

$$\frac{\partial \text{loss}^t(W)}{\partial W_{ji}} = \delta_j^t \cdot x_i^t$$

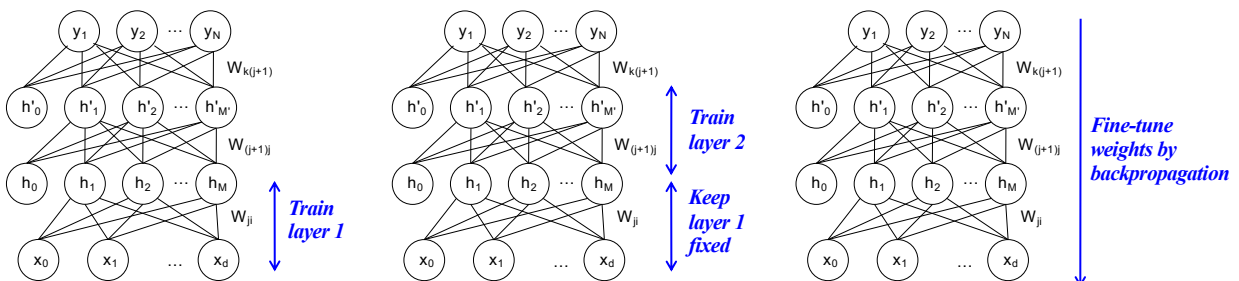
$$\delta_j^t = \sum_{(j+1)} \frac{\partial \text{loss}^t(W)}{\partial \text{net}_{(j+1)}^t} \cdot \frac{\partial \text{net}_{(j+1)}^t}{\partial \text{net}_j^t}$$

$$\delta_j^t = \left[\sum_{(j+1)} \delta_{(j+1)}^t \cdot W_{(j+1)j} \right] \cdot \sigma'(\text{net}_j^t)$$

5

Layerwise pretraining

- First, train one layer at a time, optimizing $P(x)$
- Then, fine-tune weights, optimizing $P(y|x)$ by backpropagation



6

Restricted Boltzmann machines (RBMs)

- RBM is a simple energy-based model

$$p(x, h) = \frac{1}{Z_\theta} \exp(-E_\theta(x, h))$$

$$E_\theta(x, h) = -x^T W h - b^T x - d^T h$$

It only allows h-x interactions

$$Z_\theta = \sum_{(x, h)} \exp(-E_\theta(x, h)) \leftarrow \text{normalizer}$$

- Train an RBM optimizing $P(x)$

Maximize the log-likelihood of data

$$\begin{aligned} \underbrace{\partial_{W_{ji}} \log P_W(x = x^t)}_{\text{Derivative of the log-likelihood}} &= \partial_{W_{ji}} \log \sum_h P_W(x = x^t, h) \\ &= -\partial_{W_{ji}} \log Z_W + \partial_{W_{ji}} \log \sum_h \exp(-E_W(x^t, h)) \\ &= \underbrace{-E_{p(x, h)}[x_i h_j]}_{\text{Negative phase comes from the model}} + \underbrace{E_{p(h|x=x^t)}[x_i^t h_j]}_{\text{Positive phase comes from the data}} \end{aligned}$$

The negative phase term is expensive to calculate since it requires sampling (x, h) from the model. **Contrastive divergence** is a faster solution.

7

Autoencoders

- They learn to “compress” and “reconstruct” the input data

$$\begin{aligned} \text{Encoder: } h &= \sigma(W x + b) \\ \text{Decoder: } x' &= \sigma(W' h + d) \end{aligned}$$

- Learn the weights to minimize the reconstruction loss

$$\text{loss} = \sum_t (x^t - x'^t)^2$$

- This is the same backpropagation for a network with one hidden layer, where x^t is both input and output
- They can be stacked to form a deep neural network
 - Cheaper alternatives to RBMs

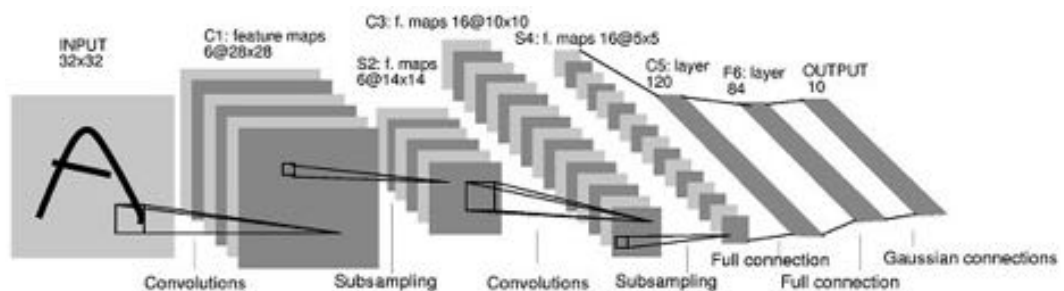
8

Basics of convolutional neural networks

9

Convolutional neural networks (CNNs)

- A CNN consists of a number of **convolutional** and **pooling (subsampling)** layers optionally followed by fully connected layers

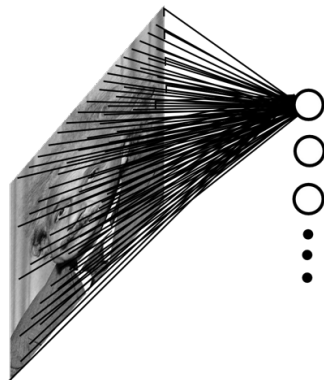


LeNet-5 by LeCun et al., 1998

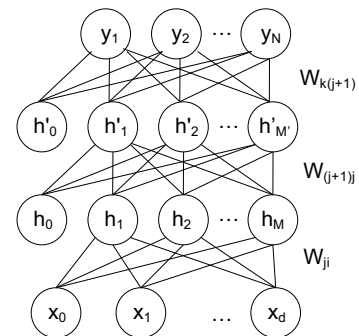
10

Fully connected layers

- When the input data is an image, a fully connected layer will produce a huge number of weights (parameters) to be learned



Example:
200x200 image
25K hidden units
→ ~1B parameters

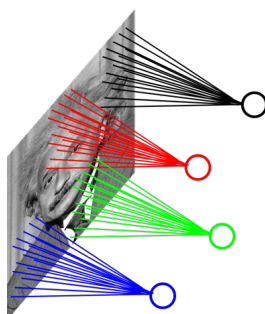


Slide credit: M.A. Ranzalo

11

Convolutional layer

- However, spatial correlation is local and statistics is similar at different locations
- Thus, small kernels are defined and their parameters are shared by all pixels
- It is convolution with learned kernels**



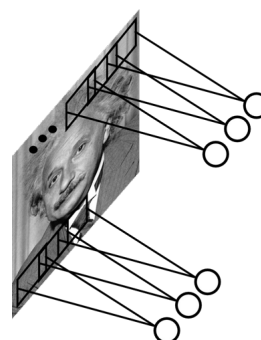
Example:
200x200 image
25K hidden units
10x10 kernels
→ ~2.5M parameters
(instead of 1B parameters)

Remember filters

Prewitt operator

$$\begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix} * \text{Image} = \text{Filtered Image}$$

Learn such filters from data

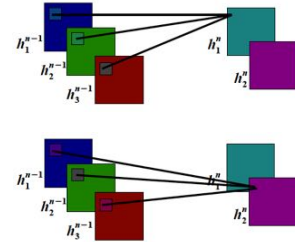
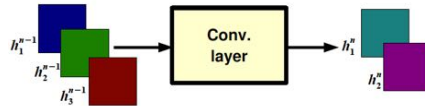
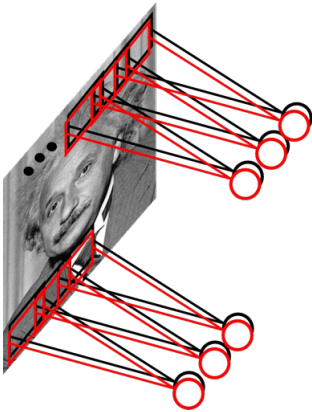


**Sparse connections
and weight sharing**

Slide credit: M.A. Ranzalo

12

Convolutional layer



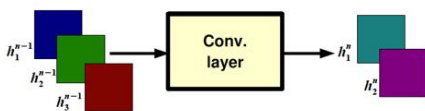
$$net_j^n = \underbrace{W_{j0}}_{\text{Bias term}} + \sum_{k=1}^K \underbrace{h_k^{n-1}}_{\substack{\text{Input feature map} \\ k^{\text{th}} \text{ feature map} \\ \text{of the previous} \\ (n-1)^{\text{th}} \text{ layer}}} * \underbrace{W_{jk}}_{\substack{\text{Kernel for the } k^{\text{th}} \\ \text{feature map of the} \\ \text{previous layer to} \\ \text{the } j^{\text{th}} \text{ feature map} \\ \text{of the current layer}}}$$

$$\underbrace{h_j^n}_{\substack{\text{Output feature map} \\ j^{\text{th}} \text{ feature map calculated for the current } n^{\text{th}} \text{ layer}}} = \text{non-linear-function}(net_j^n)$$

Slide credit: M.A. Ranzalo

13

Convolutional layer

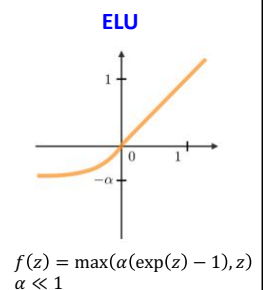
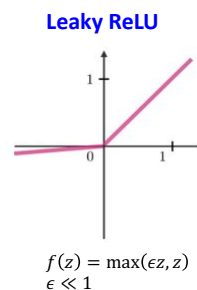
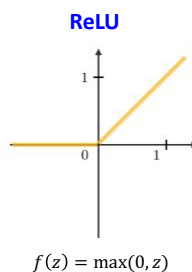


- The rectified linear unit (ReLU) is a commonly used activation function that provides nonlinearity
 - Fast to compute
 - Reduces the likelihood of the gradient to vanish
 - Better sparsity

$$net_j^n = W_{j0} + \sum_{k=1}^K h_k^{n-1} * W_{jk}$$

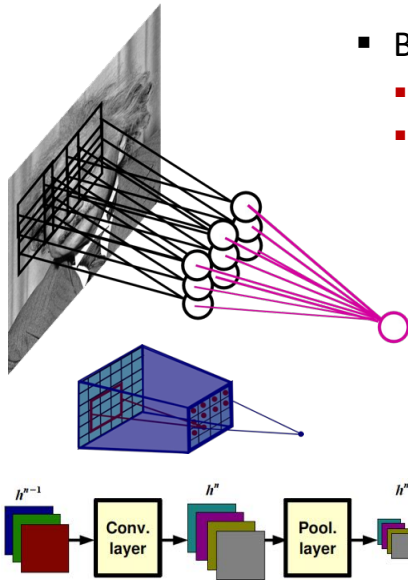
$$h_j^n = \text{non-linear-function}(net_j^n)$$

Choosing the architecture (the number of feature maps, size of kernels, and number of convolutional layers) is task dependent



14

Pooling layer



- By pooling the filter responses at different locations
 - We gain robustness to the exact location of features
 - Receptive field becomes larger for the next layer (the next layer will look at a larger spatial region)

Max-pooling:

$$h_j^n(x, y) = \max_{\substack{\bar{x} \in N(x) \\ \bar{y} \in N(y)}} h_j^{n-1}(\bar{x}, \bar{y})$$

Average-pooling:

$$h_j^n(x, y) = 1/K \sum_{\substack{\bar{x} \in N(x) \\ \bar{y} \in N(y)}} h_j^{n-1}(\bar{x}, \bar{y})$$

L2-pooling:

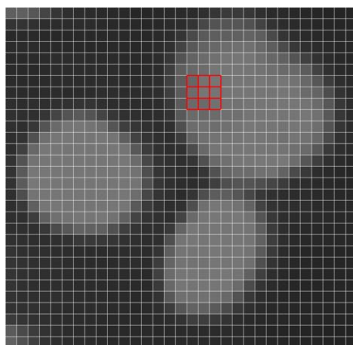
$$h_j^n(x, y) = \sqrt{\sum_{\substack{\bar{x} \in N(x) \\ \bar{y} \in N(y)}} h_j^{n-1}(\bar{x}, \bar{y})^2}$$

Slide credit: M.A. Ranzalo

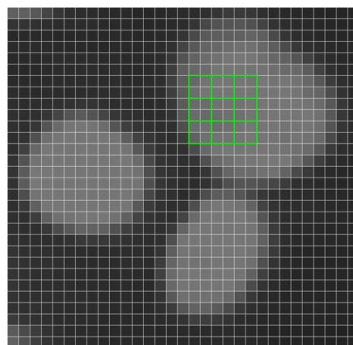
15

Receptive field

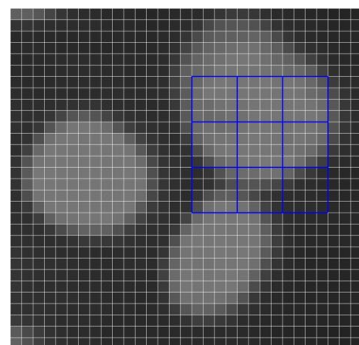
- It is the region in the input space that a convolution can “see”



What a 3x3 convolutional filter sees



What the 3x3 convolutional filter sees after 2x2 pooling

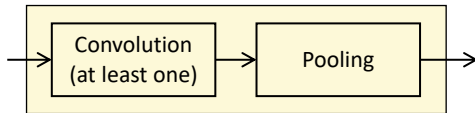


What the 3x3 convolutional filter sees after applying 2x2 pooling twice

16

Typical design

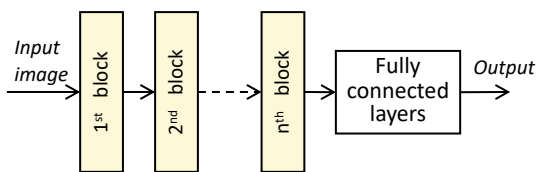
One block



After one block

- Number of feature maps is usually increased (conv. layer)
- Spatial resolution is usually decreased (pooling layer and/or stride in conv. layer)
- Receptive field gets larger

Whole system



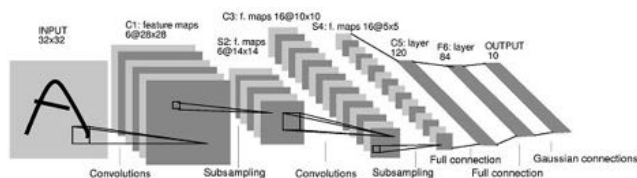
After several blocks

- Spatial resolution is greatly reduced and number of feature maps is large so convolution would not make any sense
- Next layer(s) will consist of fully connected layers (with or without hidden layers)
- Sigmoid and softmax layer is used at the end for binary and multi-class classification, respectively

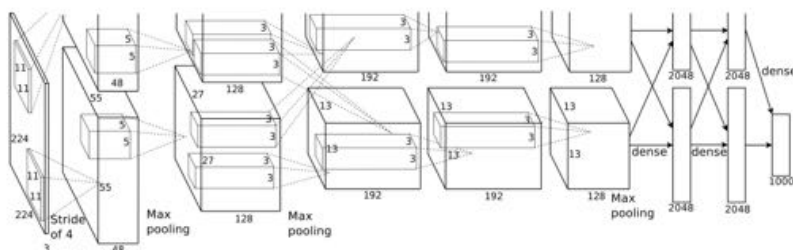
All layers are differentiable so that standard backpropagation can be used

17

Example CNN architectures



LeNet-5 by LeCun et al., 1998

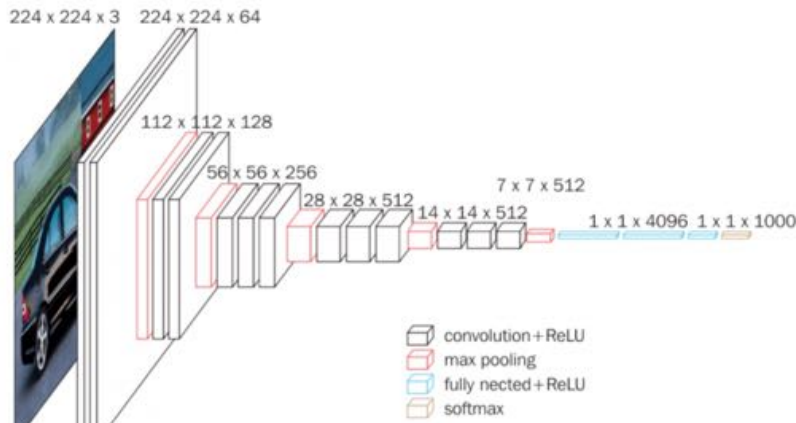


AlexNet by Krizhevsky et al., 2012. ImageNet classification with deep convolutional neural networks.

https://www.cs.toronto.edu/~kriz/imagenet_classification_with_deep_convolutional.pdf

18

Example CNN architectures



VGG16 by Simonyan and Zisserman, 2015. Very deep convolutional networks for large-scale image recognition.
<https://arxiv.org/pdf/1409.1556.pdf>

19

Example CNN architectures

type	patch size/ stride	output size	depth	# 1x1	# 3x3 reduce	# 3x3	# 5x5 reduce	# 5x5	pool proj	params	ops
convolution	7x7/2	112x112x64	1							2.7K	34M
max pool	3x3/2	56x56x64	0								
convolution	3x3/1	56x56x192	2		64	192				112K	360M
max pool	3x3/2	28x28x192	0								
inception (3a)		28x28x256	2	64	96	128	16	32	32	159K	128M
inception (3b)		28x28x480	2	128	128	192	32	96	64	380K	304M
max pool	3x3/2	14x14x480	0								
inception (4a)		14x14x512	2	192	96	208	16	48	64	364K	73M
inception (4b)		14x14x512	2	160	112	224	24	64	64	437K	88M
inception (4c)		14x14x512	2	128	128	256	24	64	64	463K	100M
inception (4d)		14x14x528	2	112	144	288	32	64	64	580K	119M
inception (4e)		14x14x832	2	256	160	320	32	128	128	840K	170M
max pool	3x3/2	7x7x832	0								
inception (5a)		7x7x832	2	256	160	320	32	128	128	1072K	54M
inception (5b)		7x7x1024	2	384	192	384	48	128	128	1388K	71M
avg pool	7x7/1	1x1x1024	0								
dropout (40%)		1x1x1024	0								
linear		1x1x1000	1							1000K	1M
softmax		1x1x1000	0								

Table 1: GoogLeNet incarnation of the Inception architecture

GoogLeNet by Szegedy et al., 2014. Going deeper with convolutions.
<https://arxiv.org/abs/1409.4842>

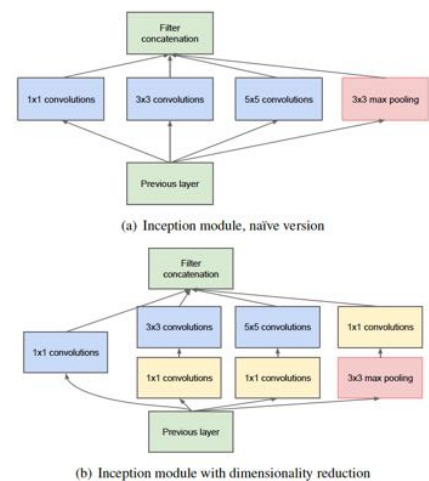


Figure 2: Inception module

20

Example CNN architectures

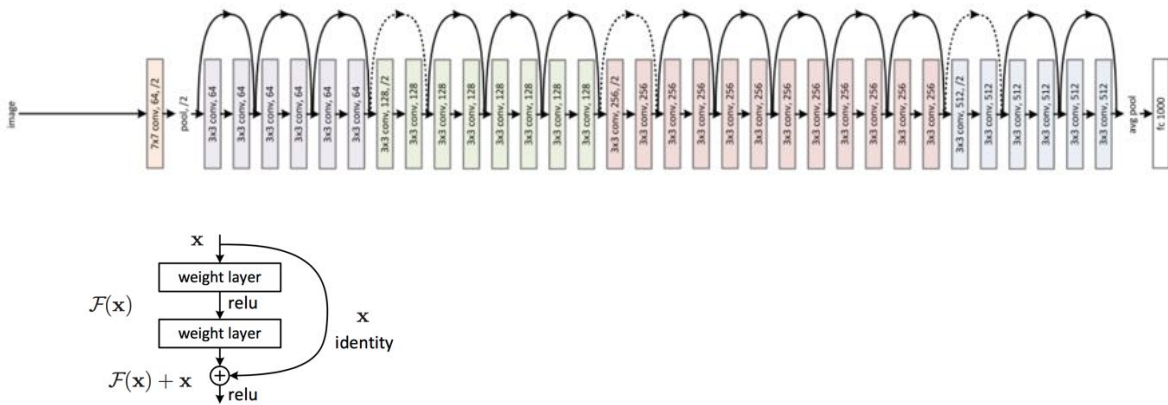


Figure 2. Residual learning: a building block.

ResNet by He et al., 2015. Deep residual learning for image recognition.
<https://arxiv.org/pdf/1512.03385.pdf>

21

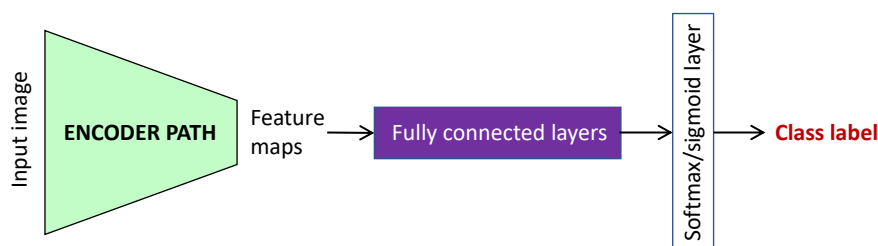
CNNs for medical image classification

Please check the following survey paper for more about deep learning in medical images and more references.
 Litjens et al., *A survey on deep learning in medical image analysis*, *Medical Image Analysis*, 2017.
<https://www.sciencedirect.com/science/article/pii/S1361841517301135>

22

CNNs for image classification

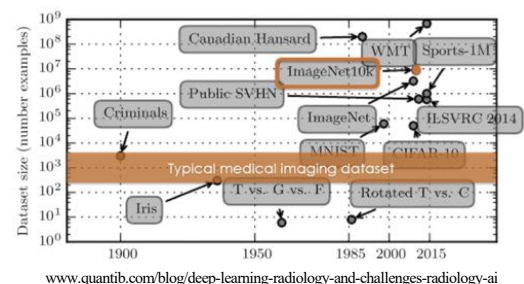
- A CNN compresses an image into a set of feature maps to capture semantic/contextual information from the image
- This compression corresponds to downsampling the image using convolution and pooling layers
- Then it puts fully connected layers on the top of the feature maps to predict a class for the entire image



23

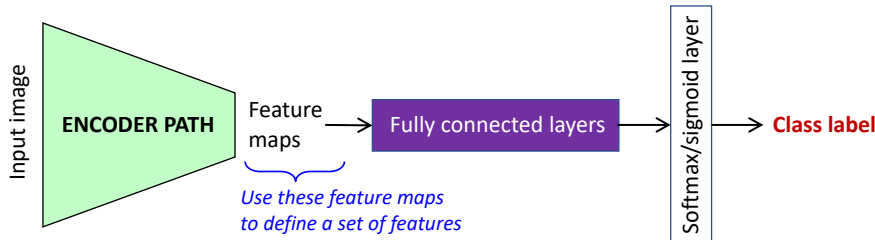
Transfer learning

- Dataset sizes are typically small
- Thus, it is popular to use transfer learning, which employs networks (and thus, their learned weights) previously trained on large datasets
- Two main approaches
 1. Use a pretrained network as a feature extractor
 2. Finetune a pretrained network on the medical data

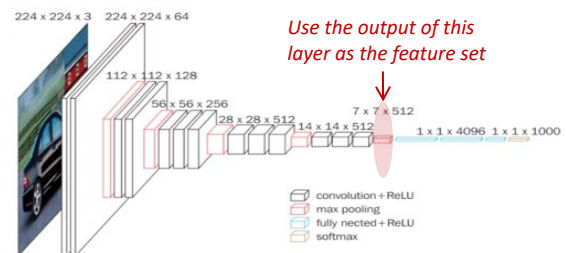


24

Use a pretrained network as a feature extractor

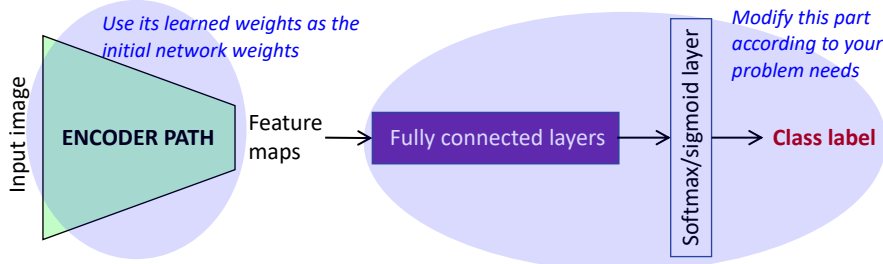


1. Feed your image to the pretrained network (may need to resize the image)
2. Calculate feature maps using the learned weights
3. Use the feature maps at the last layer of the encoder to define a feature set for your image (also possible to use feature maps of the other layers)
4. Use a classifier on this feature set

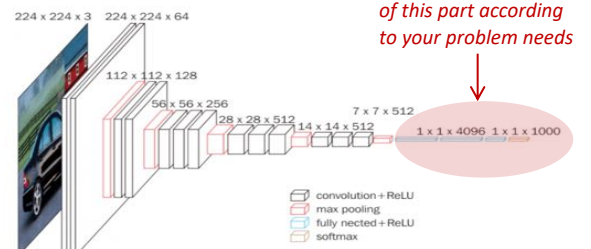


25

Finetune a pretrained network on the medical data

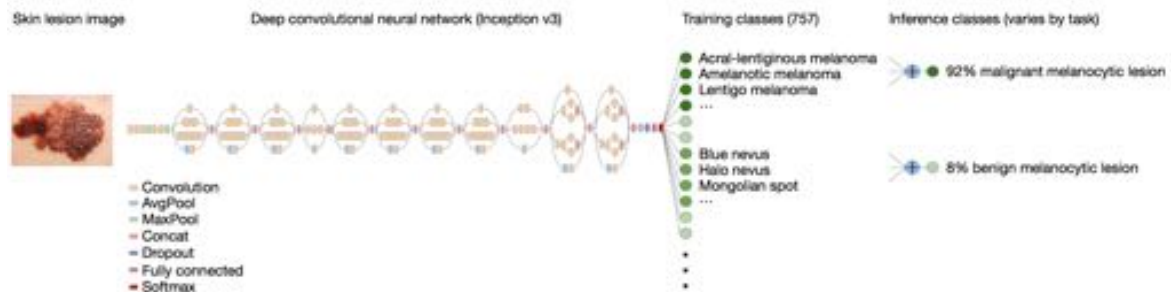


1. Feed your image to the pretrained network (may need to resize the image)
2. Modify the network architecture after the encoder path. Changing the last softmax/sigmoid layer according to your classification problem is a must. May also need to change the fully connected layers.
3. Use the learned weights (of the encoder) as the initial network's weights
4. Finetune the weights by backpropagation on your own medical data



26

Example: CNN for skin cancer classification

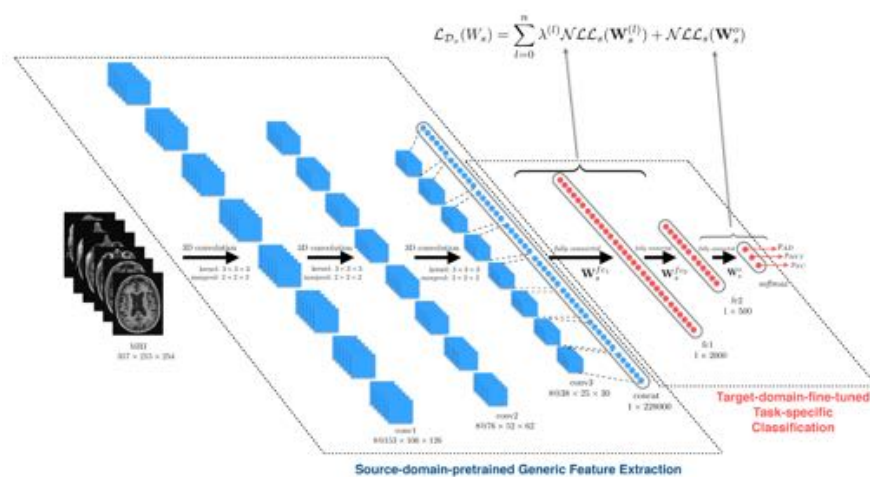


The authors used the Google Inception v3 CNN architecture pretrained on the ImageNet dataset (1.28 million images over 1,000 generic object classes) and finetuned on their own dataset of 129,450 skin lesions. They resized each image to 299x299 pixels to make it compatible with the original dimensions of the Inception v3 network architecture. They defined 757 training classes, for which the probabilities would be accumulated to infer the final inference class.

Esteva et al., 2017. Dermatologist-level classification of skin cancer with deep neural networks. *Nature* 542, 115–118.
<https://www.nature.com/articles/nature21056>

27

Example: 3D CNN for Alzheimer's disease diagnostics



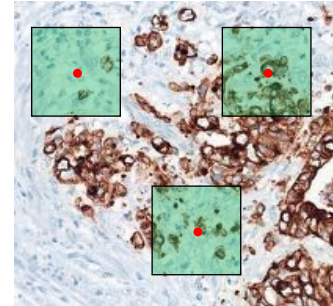
Hosseini-Asl et al., 2016. Alzheimer's disease diagnostics by a deeply supervised adaptable 3D convolutional network.
<https://arxiv.org/abs/1607.00556>

28

CNNs for object detection and segmentation

Training:

- Small patches are cropped around individual pixels
- Each patch is labeled with the class of the pixel, around which it is cropped
- CNN is trained on these small patches



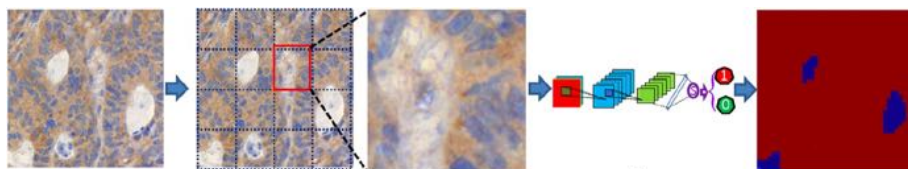
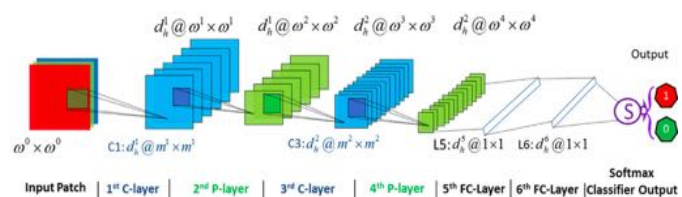
Detection/segmentation:

- For an entire (large) image, patches are obtained using a sliding window approach
- These patches are classified by the trained CNN
- Outputs (i.e., posteriors) generated by this CNN are commonly postprocessed

29

Example: CNN for histopathological image segmentation

- *Slide a window* over an image to obtain patches
- Using the trained CNN, classify each patch with either the epithelial or the stromal class



Xu et al., 2016. A deep convolutional neural network for segmenting and classifying epithelial and stromal regions in histopathological images. *Neurocomputing* (191), 214-223.

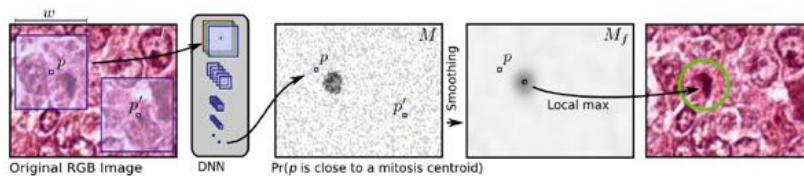
<https://www.sciencedirect.com/science/article/pii/S0925231216001004>

30

Example: CNN for mitosis detection

- **Slide a window** over an image to obtain patches
- Using the trained CNN, obtain the probability of a pixel belonging to a mitotic cell
- Find the local maxima on the smoothed probability map

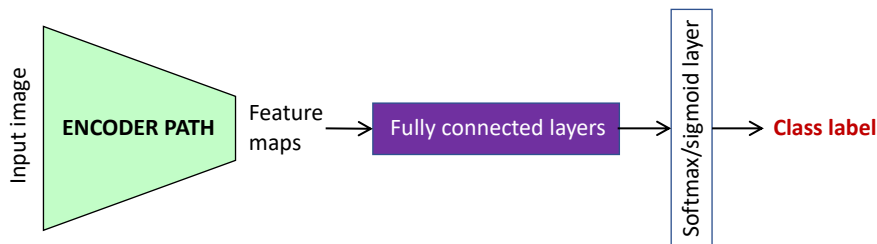
1. *Classifying each pixel in a sliding window fashion, used by earlier studies, is expensive as it requires lots of redundant calculations*
 2. *Trade-off between localization accuracy and the use of context*
 - *Larger patches require more max-pooling layers that reduce the localization accuracy*
 - *Small patches results in seeing only little context*
- Dense prediction networks, used by recent studies, have greatly improved efficiency and accuracy.*



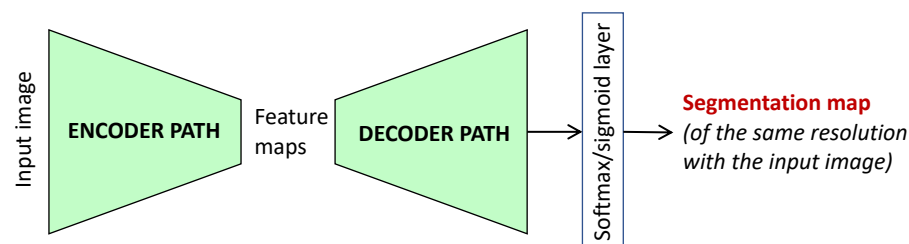
Ciresan et al., 2013. Mitosis detection in breast cancer histology images with deep neural networks. MICCAI.
https://link.springer.com/chapter/10.1007/978-3-642-40763-5_51

31

Convolutional neural networks (CNNs) for image classification



Dense prediction networks for semantic image segmentation



32

Thank you!

Next time:

Dense prediction networks